


# Chapter 2

## Application Layer



Computer Networking: A Top-Down Approach  
8th Edition, 2020, Pearson,  
James F. Kurose, Keith W. Ross

# Chapter 2 outline

## 2.1 Principles of network applications

## 2.2 Web and HTTP

## 2.3 Electronic mail

- SMTP, POP3, IMAP

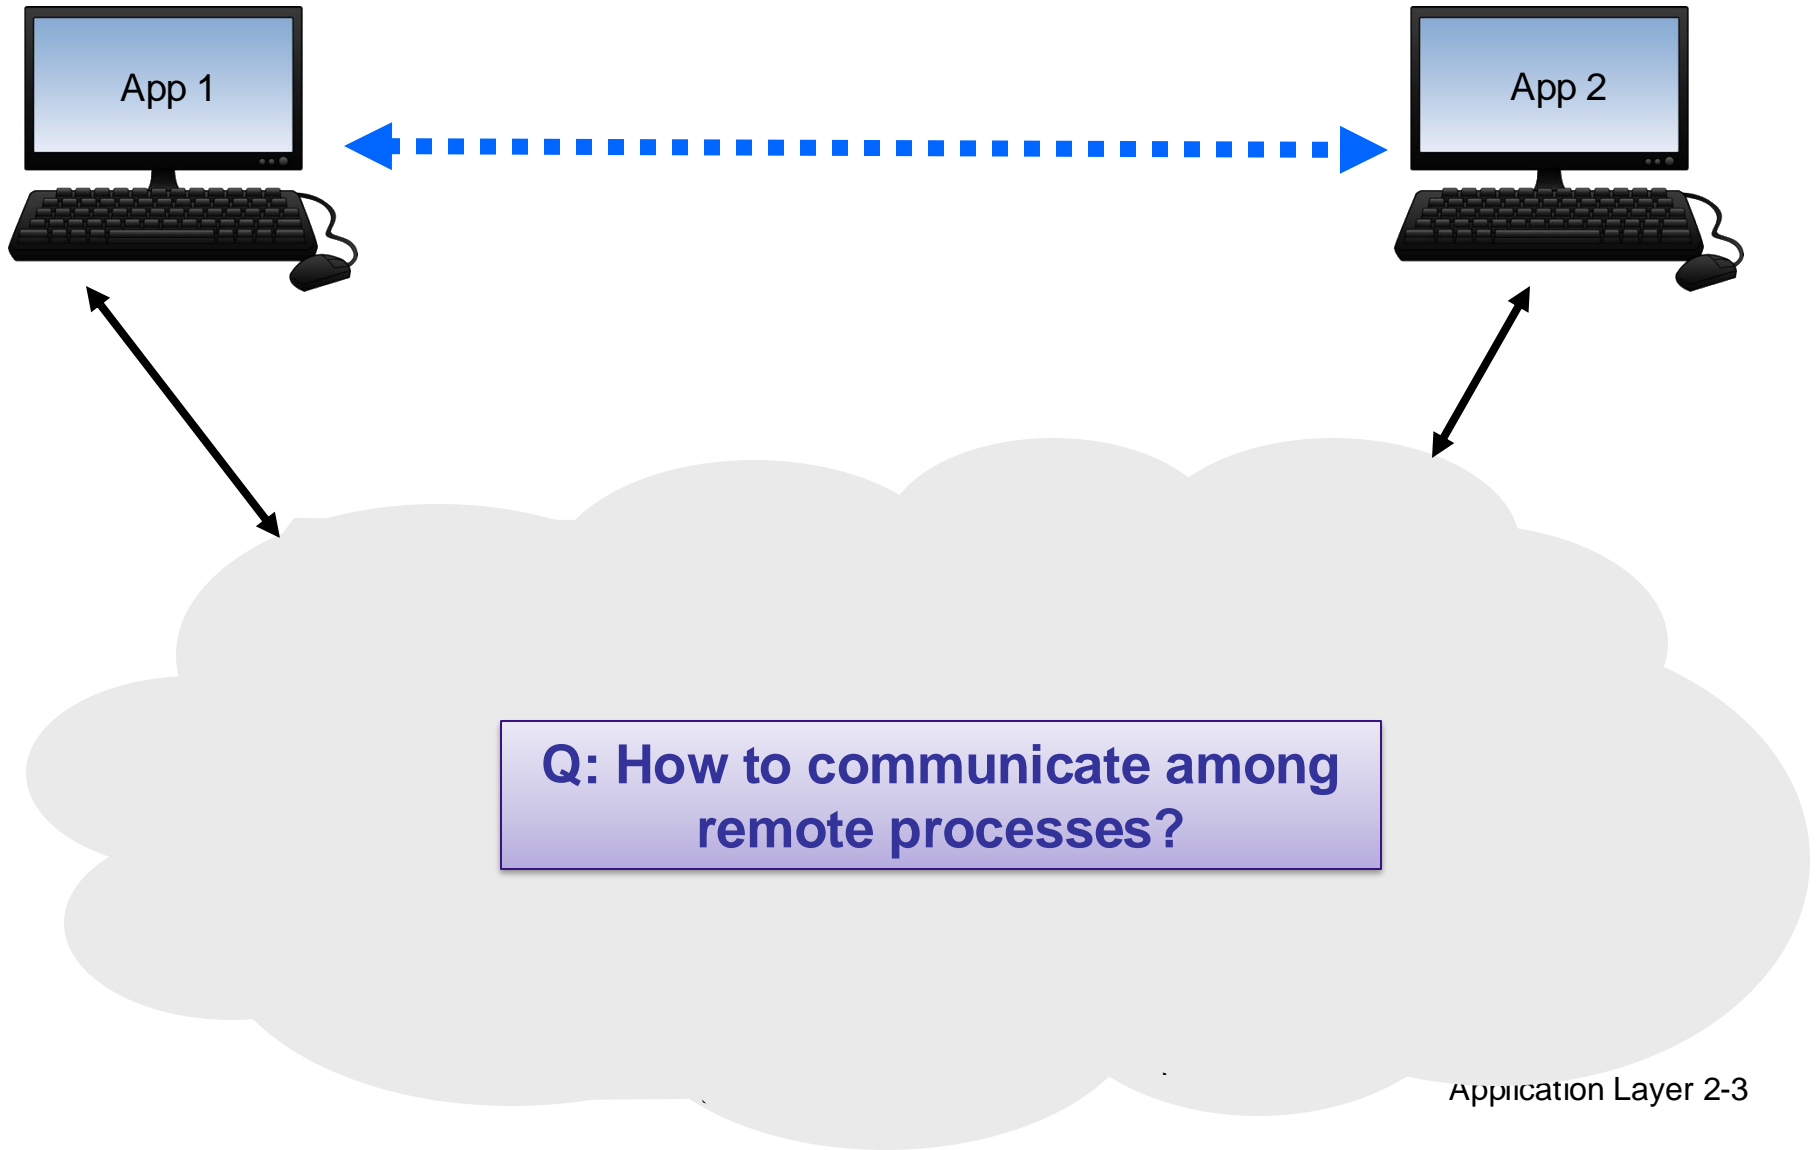
## 2.4 DNS

## 2.5 P2P applications

[2.6 video streaming and content distribution networks]

## 2.7 Socket programming with UDP and TCP

# Focus : Application Layer



# Network Applications

- e-mail
- web
- text messaging
- remote login
- P2P file sharing
- multi-user network games
- streaming stored video (YouTube, Hulu, Netflix)
- voice over IP (e.g., Skype)
- real-time video conferencing
- social networking
- search



**Q: Which application architectures can we identify?**


# Client-Server Architecture

## server:

- always-on host
- permanent IP address
- data centers for scaling

## clients:

- communicate with server
- may be intermittently connected
- may have dynamic IP addresses
- do not communicate directly with each other



# Application Client - Server

## CLIENT : “*active open*”


- invoked directly by user
- local on user's personal computer
- actively initiates contact with server
- one session at a time (per (instance of) application)
- access multiple application(s) (instances) as needed
- simple hardware and software

e.g. : Thunderbird, Outlook, Chrome, Firefox, Safari, ...

## SERVER : “*passive open*”

- special purpose program for one application
- can handle multiple remote clients at same time
- runs on a shared computer
- waits passively for contact from arbitrary remote client
- publicly known address, port
- powerful hardware and sophisticated operating system
- server program also called daemon (e.g. FTP daemon)


e.g.: Apache, IIS, Sendmail, Postfix, ...



Information can flow in both directions between client and server  
Some application programs act as client and server

# P2P Architecture

- **no** always-on server
- arbitrary end systems directly communicate
- peers request service from other peers, provide service in return to other peers
  - **self scalability** – new peers bring new service capacity, as well as new service demands
- peers are intermittently connected and change IP addresses
  - complex management

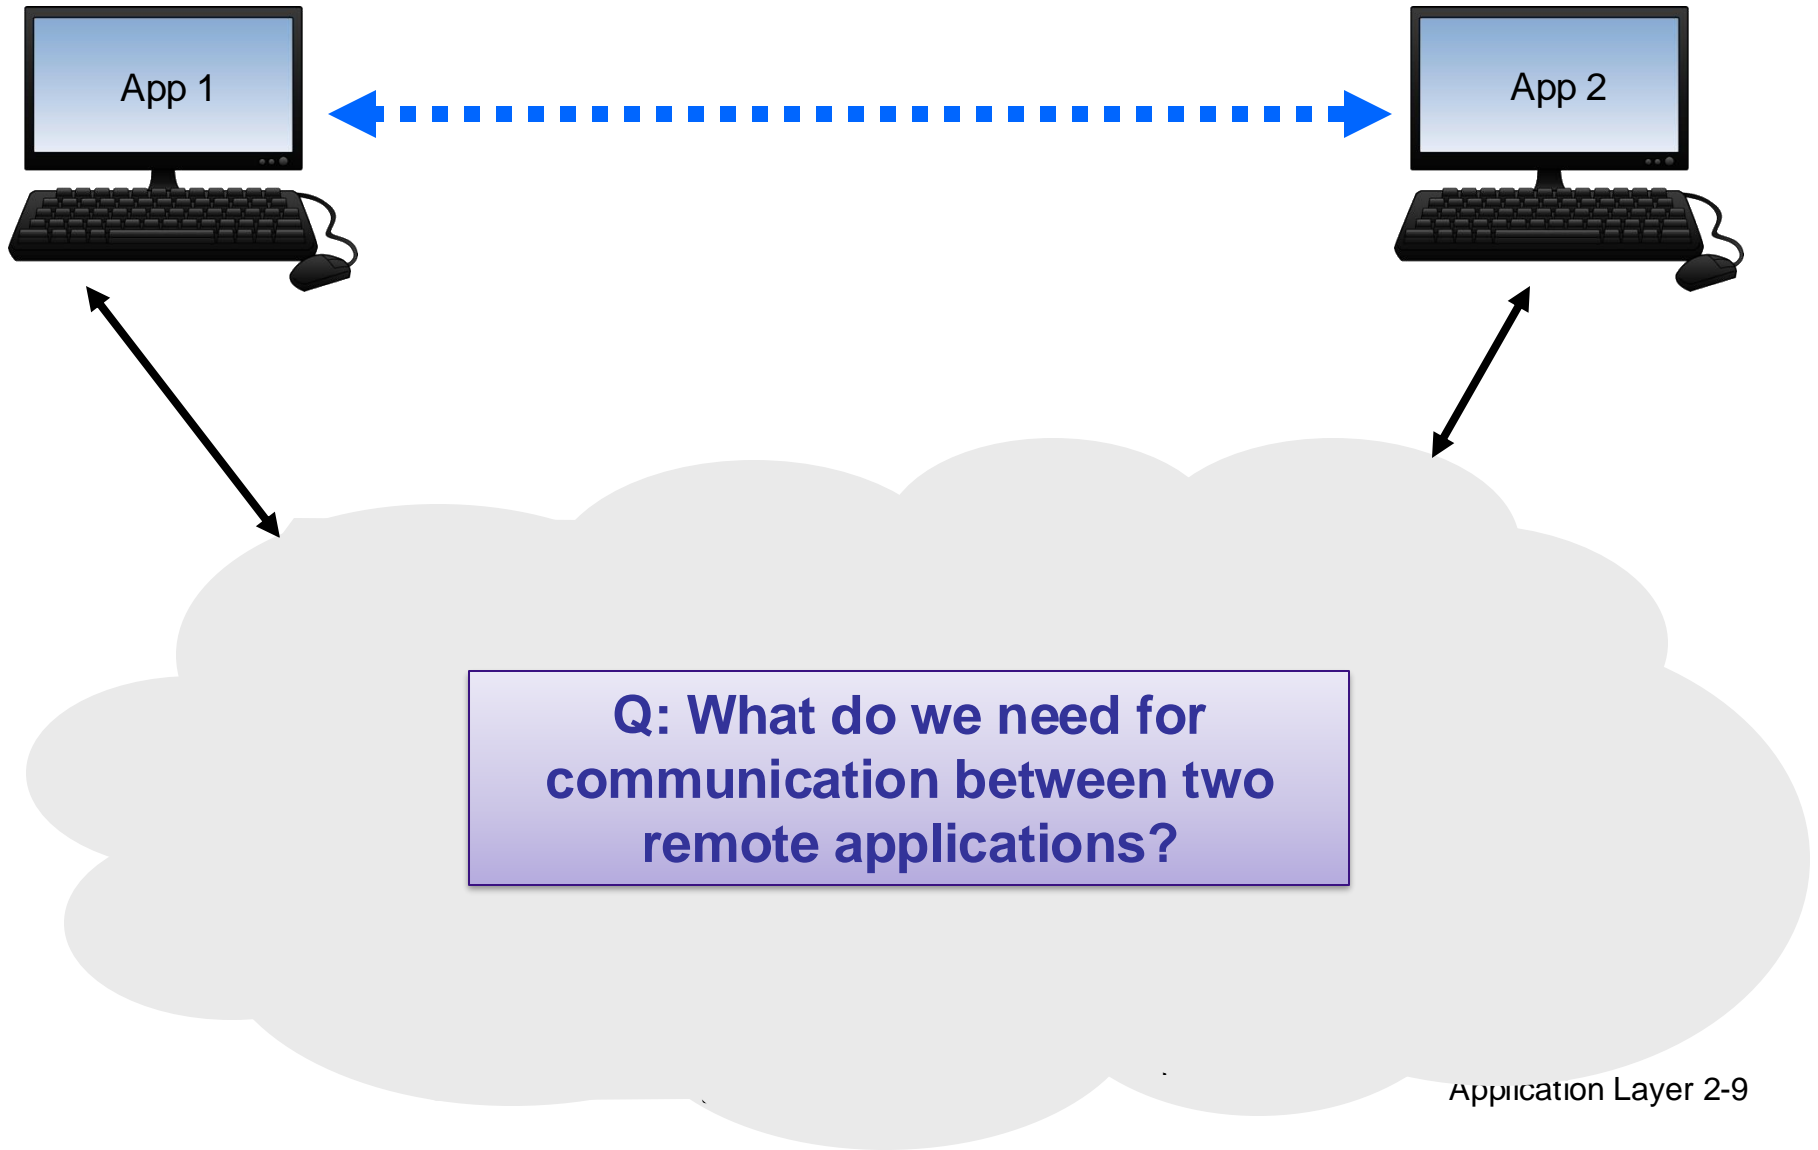


# Hybrid of client-server and P2P

- Skype
  - voice-over-IP P2P application
  - centralized server: finding address of remote party
  - client-client connection: direct (not through server)
- Instant messaging
  - chatting between two users is P2P
  - centralized service: client presence detection/location
    - user registers its IP address with central server when it comes online
    - user contacts central server to find IP addresses of contacts



# Focus : Application Layer



# What is needed for application-layer communication?



## 1. Addressing of remote processes

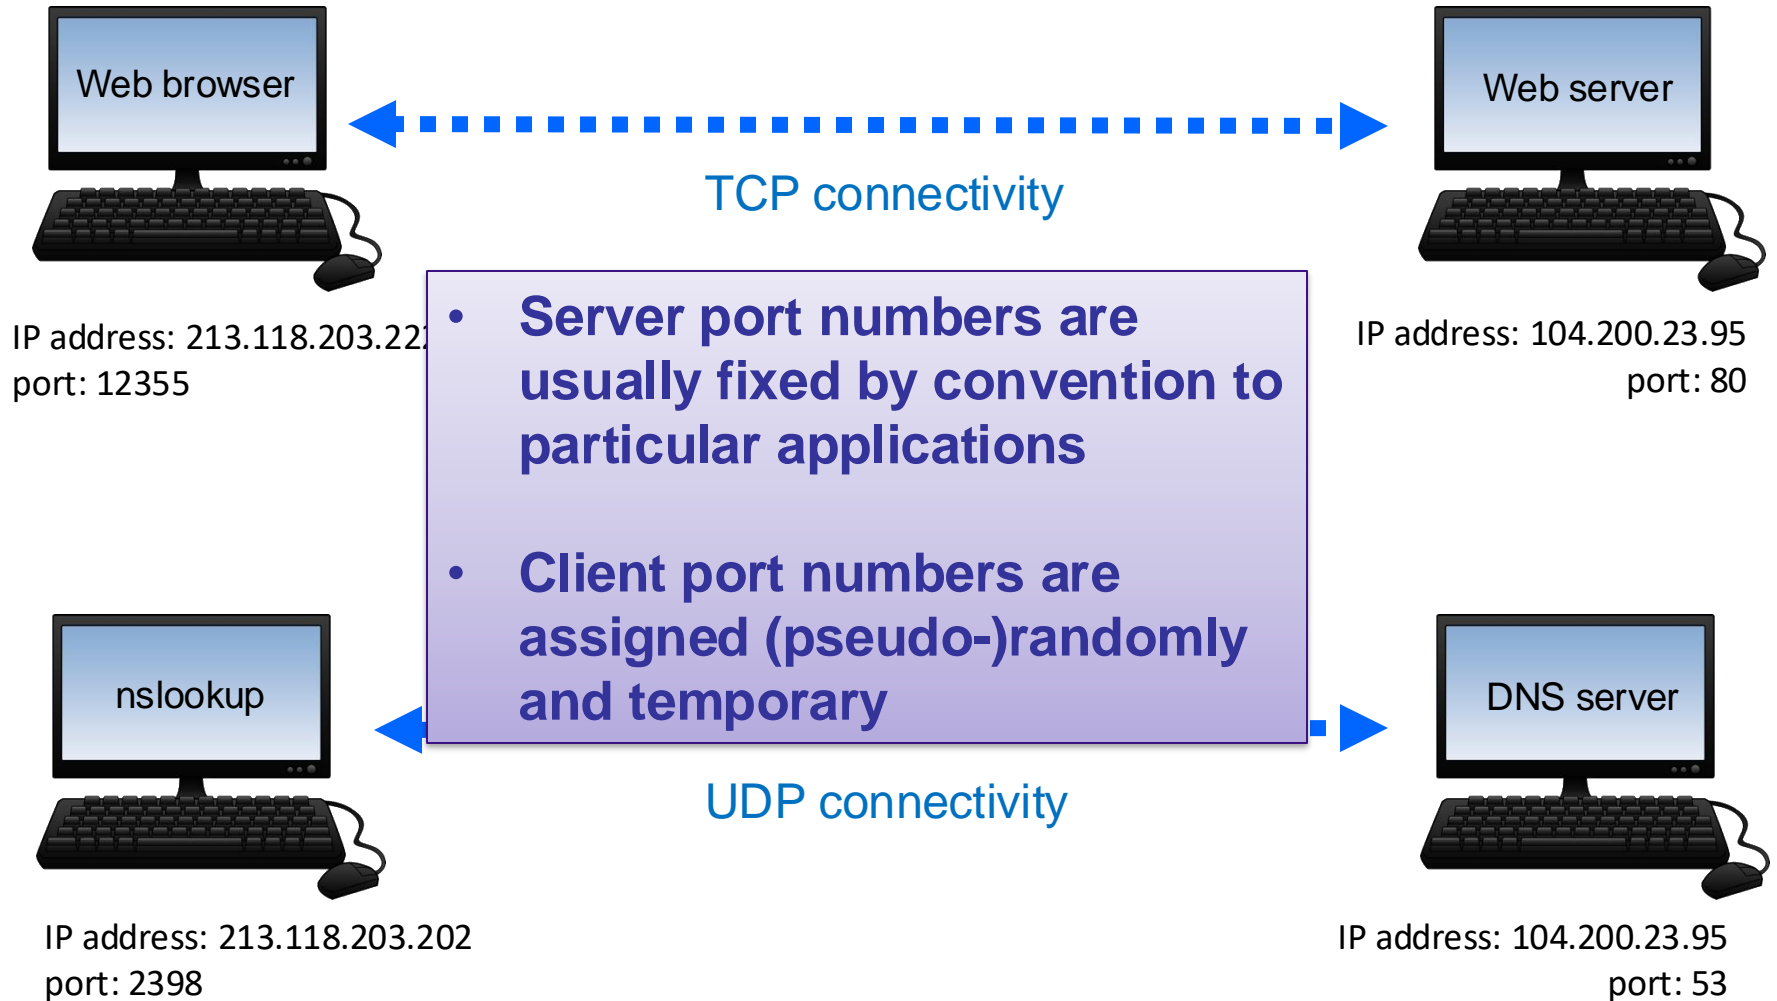
- Nodes (i.e., interface) on the Internet are addressed using IP addresses (32-bit IPv4 address or 128-bit IPv6 address)
- Applications (i.e., OS processes) on the considered nodes are addressed using a port number (16-bit value, i.e. 0-65535)

## 2. Type of connectivity

- Provided by the lower layer = transport layer protocol
- Two (most used) types:
  - Reliable byte stream service = TCP
  - Best-effort datagram service = UDP

# Examples: web browsing & DNS

## 1. Addressing



# Common server port mappings for applications

Application	Usual server port	Transport Layer Protocol
FTP	20,21	TCP
SSH	22	TCP
Telnet	23	TCP
SMTP (email transfer server)	25	TCP
DNS	53	UDP/TCP
HTTP (web)	80	TCP
POP (email server)	109, 110	TCP

Note that these are assigned by convention (reserved as per RFC 1700).  
Port numbers > 1024 can be freely used (e.g., web server at port 443)

```
$ sudo netstat -plnt
```

```
Active Internet connections (only servers)
```

Proto	Recv-Q	Send-Q	Local Address	Foreign Address	State	PID/Program name
tcp	0	0	0.0.0.0:3306	0.0.0.0:*	LISTEN	3686/mysqld
tcp	0	0	:::443	:::*	LISTEN	2218/httpd
tcp	0	0	:::80	:::*	LISTEN	2218/httpd
tcp	0	0	:::22	:::*	LISTEN	1051/sshd

# What Transport Service Does An App Need?

## data integrity

- some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
- other apps (e.g., audio) can tolerate some loss

## timing

- some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

## throughput


- some apps (e.g., multimedia) require minimum amount of throughput to be “effective”
- other apps (“elastic apps”) make use of whatever throughput they get

## security


- encryption, data integrity, ...

# Type of connectivity (transport layer)

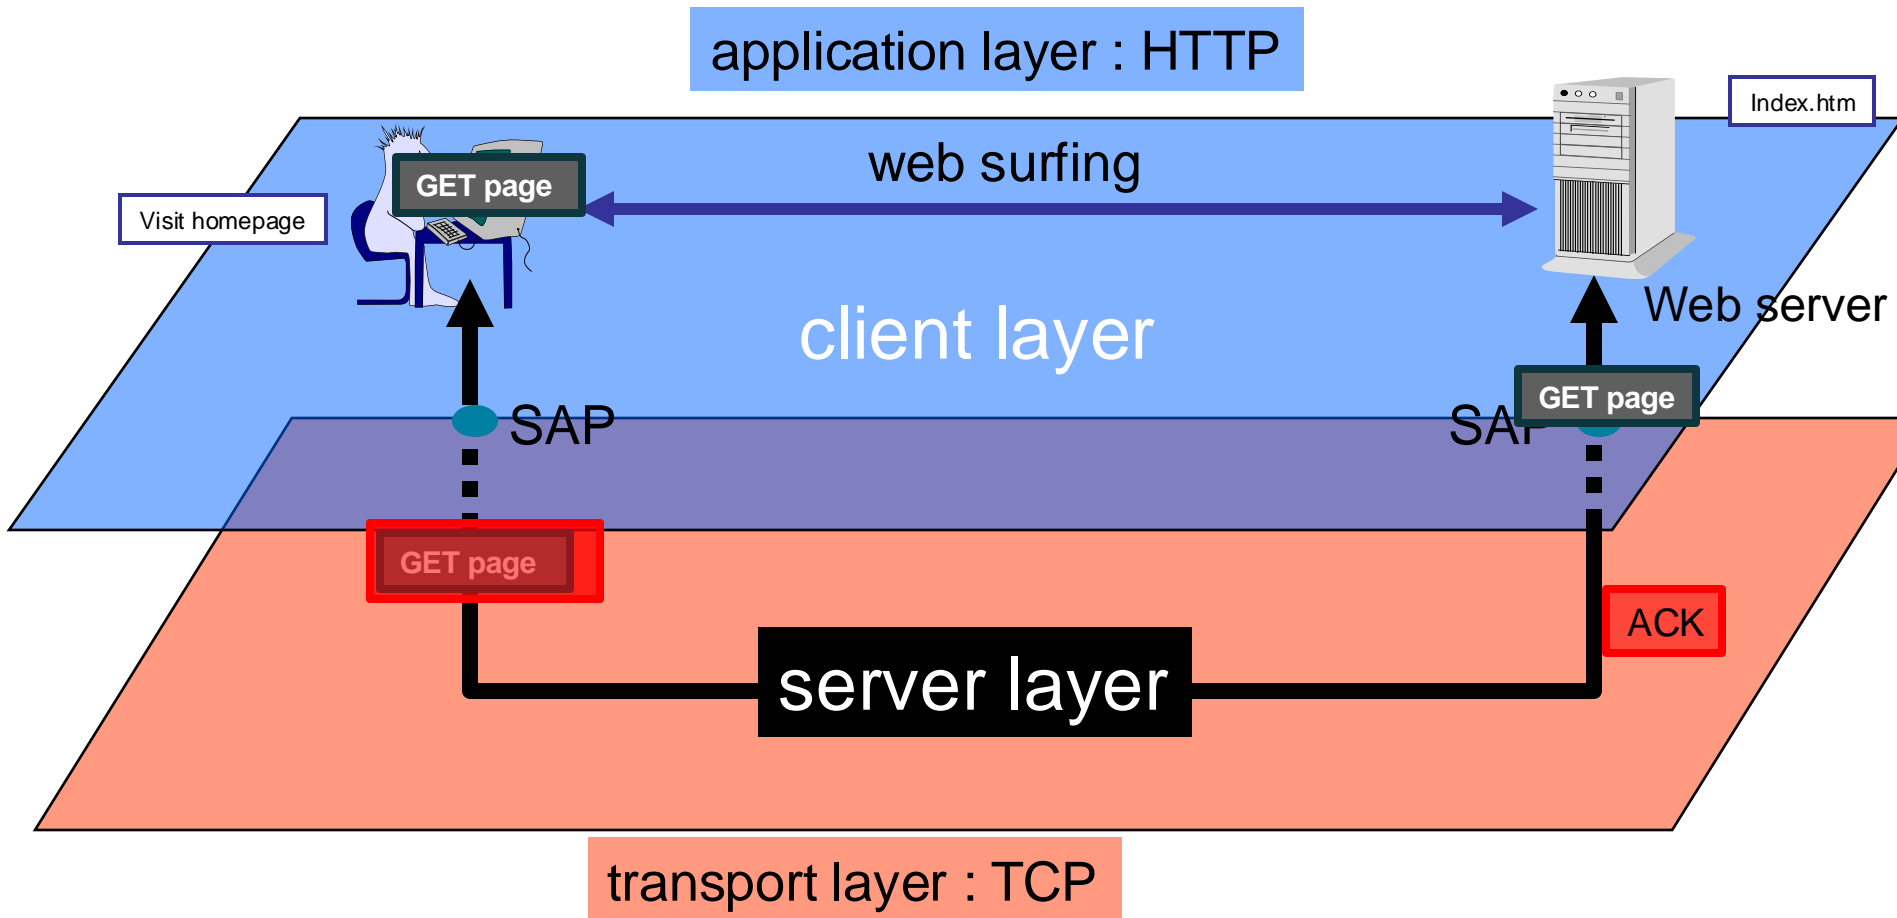
- TCP = reliable byte stream service
  - Connection setup before communication (3-way handshake)
  - Ensures every byte has arrived in order (acknowledgements)



- UDP = best-effort datagram service
  - UDP treats them as separate messages (of max 65507 bytes)
  - Application itself must handle lost messages



# Application/Transport Layer Interaction

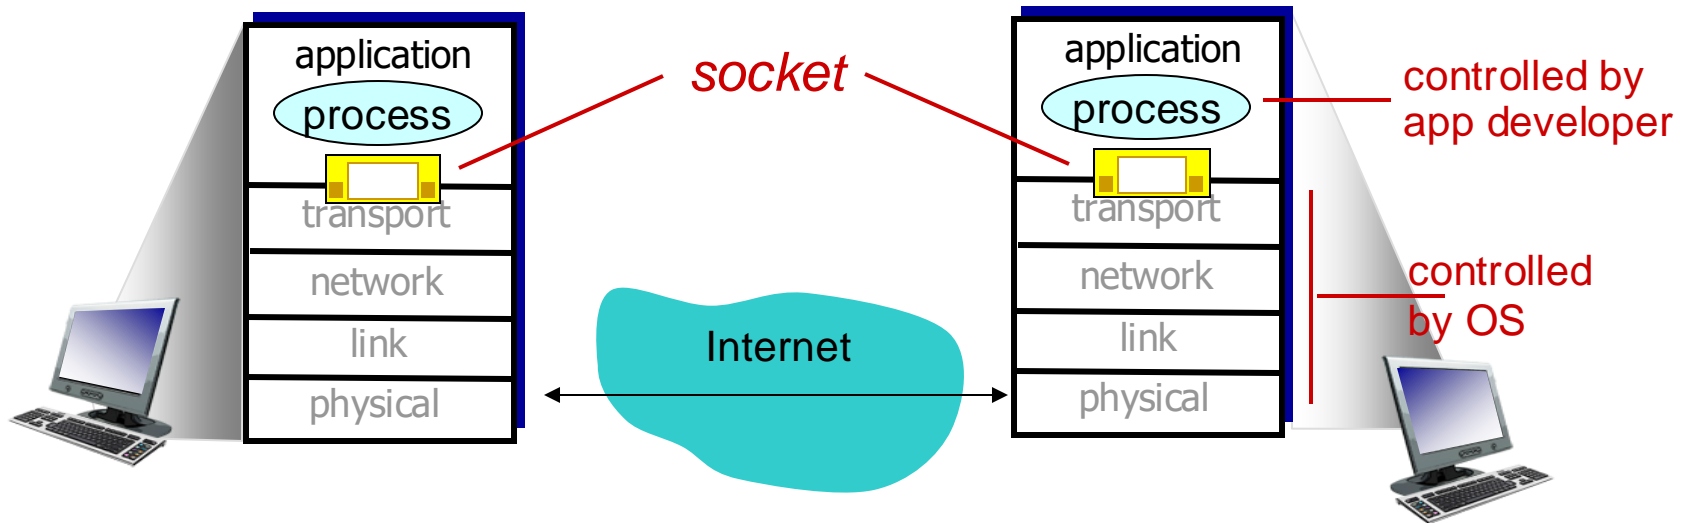


SAP : Service Access point

Remark : client - server

# Network programming using sockets

- Network communication from one process to another remote process relies on **Operating System Calls for network programming**
- Most common API is based on **Berkeley Sockets**:
  - Socket = communication endpoint, referring to the “door” between the application process and the end-to-end transport protocol





# Socket Example (in Python 3)

## TCP Hello Server

```
from socket import *

s = socket(AF_INET, SOCK_STREAM)
s.bind(('127.0.0.1', 9999))
s.listen(5)

while True:
    c, addr = s.accept()
    g = bytes("Hello %s" % a[0])
    c.send(bytes(g, 'utf-8'))
    c.close()
```

# start server first

```
$ python tcp-server.py
```

## TCP Hello Client

```
from socket import *

s = socket(AF_INET, SOCK_STREAM)
s.connect(('127.0.0.1', 9999))

s.send(b'Hello, world')
data = s.recv(1024)

s.close()

print 'Received', data
```

# next start client

```
$ python tcp-client.py
Received b'Hello 127.0.0.1\n'
```

# Applications and application-layer protocols



- Network Application : communicating, distributed processes (e.g. Web, e-mail, P2P file sharing, ...)
  - Applications are responsible for WHAT is sent between the processes

**Q: How to enable different user interfaces for the same application (e.g., Chrome vs. Firefox, Outlook vs. Thunderbird)?**

- Applications use an **application-layer protocol** to implement these services
- The **application layer** is the collection of these protocols (FTP, SMTP, POP, IMAP, HTTP, DNS, SNMP, etc.)
  - **Example:** HTTP protocol: GET and POST messages

# App-layer protocol defines

- **types of messages exchanged,**  
e.g., request, response
- **message syntax:**  
what fields in messages &  
how fields are delineated
- **message semantics**  
meaning of information in  
fields
- **rules** for when and how  
processes send & respond  
to messages

## **open protocols:**

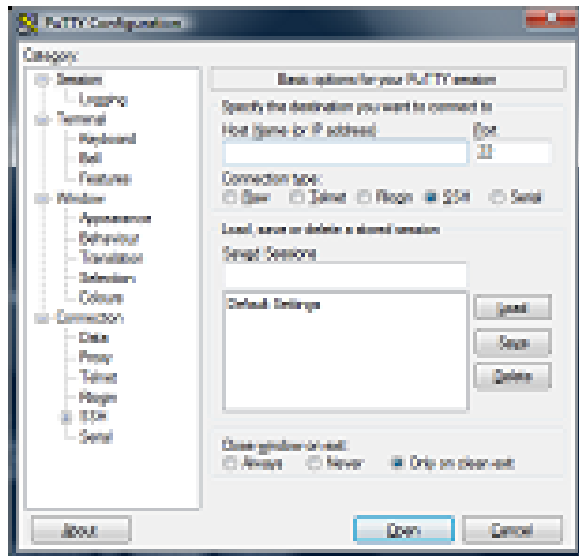
- defined in RFCs
- allows for  
interoperability
- e.g. HTTP, SMTP

## **proprietary protocols:**

- e.g. Skype, BitTorrent

# Telnet client

- You can use a **telnet** client (software) to play with the protocols: HTTP, SMTP and others.
- The **telnet** client is a 'generic' TCP client:
  - Sends whatever you type to the TCP socket.
  - Prints whatever comes back through the TCP socket
  - Useful for testing TCP servers (ASCII based protocols).



Putty (Windows)

```
handbook@trusty: ~  
handbook@trusty:~$ telnet 192.168.1.3  
Trying 192.168.1.3...  
Connected to 192.168.1.3.  
Escape character is '^]'.  
Ubuntu 14.04 LTS  
trusty login: handbook  
Password:  
Last login: Sat Dec 20 23:12:58 CST 2014 from 192.168.1.3 on pts/12  
Welcome to Ubuntu 14.04 LTS (GNU/Linux 3.13.0-24-generic x86_64)  
  
* Documentation: https://help.ubuntu.com/
```

telnet command line (Linux)

# Investigate simple protocols using Telnet

Some application-layer protocols which can be tested with the telnet client:

- ECHO (port 7)
- DAYTIME (port 13)
- TELNET (port 23)
- HTTP (port 80)
- ...

Ask creation of a socket to 157.193.40.10 at port 13

## Daytime example

```
$ telnet 157.193.40.10 13
Trying 157.193.40.10...
Connected to eduser2.ugent.be.
Mon Oct 2 14:27:43 1998
Connection closed by foreign host.
```

## Echo example

```
$ telnet 157.193.40.10 7
Trying 157.193.40.10...
Connected to eduser2.ugent.be.
Hallo
Hallo
^C
```

We type Hallo and see it on the display  
(Telnet operation)

We receive it a second time on the display  
(Echo operation)

^C will close the connection

## Telnet protocol example

```
$ telnet 157.193.40.10 23
Trying 157.193.40.10...
Connected to eduser2.ugent.be.
login:wtaverni
passwd:****
Last login: Mon 23 17:12:23 CET
2019 from 213.118.203.222 on pts/15
> ls
test.txt
```

# Chapter 2 outline

**2.1** Principles of network applications

**2.2 Web and HTTP**

**2.3** electronic mail

- SMTP, POP3, IMAP

**2.4** DNS

**2.5** P2P applications

[**2.6** video streaming and content distribution networks]

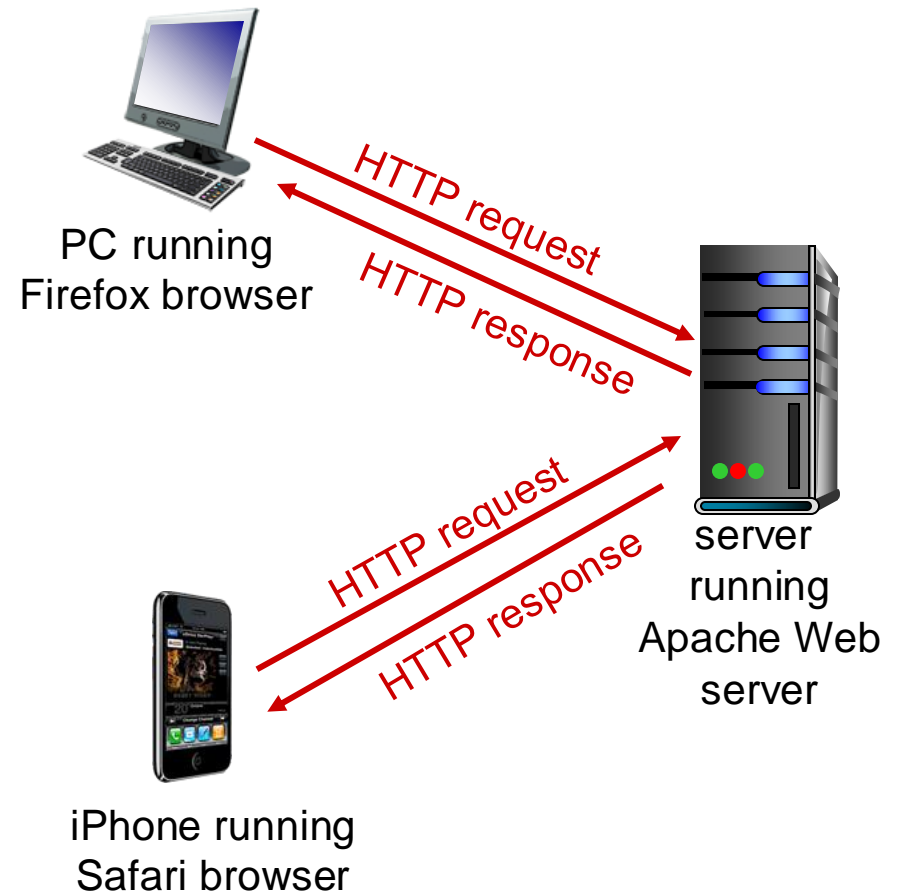
**2.7** Socket programming with UDP and TCP

# HTTP overview

## HTTP: HyperText Transfer Protocol


- Web's application layer protocol to exchange Web objects (e.g., html pages, videos, pictures)
- client/server model
  - **client**: browser that requests, receives, and “displays” Web objects
  - **server**: Web server sends objects in response to requests
- HTTP 1.0: RFC 1945
- HTTP 1.1: RFC 2068
- HTTP 2.0: RFC 7540
- HTTP 3.0: RFC 9114

text-based



# Web object example

URL :      <IP address>      : port / <filepath of object>  
         <domain/host name>



ComNet.html

Base HTML-file which includes  
several referenced objects

ietf.gif

First referenced object (with own  
url: [https://ietf.org/media/images/ietf-  
logo.original.png](https://ietf.org/media/images/ietf-logo.original.png) )

ugent.png

Second referenced object ...

IDLab.jpg

Third referenced object ...

Example weg page (Communication Networks)



I E T F®



UNIVERSITEIT  
GENT

IDLab  
INTERNET & DATA LAB



# HTTP overview

## uses TCP:

- server is listening (open socket, port 80)
- client initiates TCP connection (creates socket, port >1024) to server, port 80
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

## HTTP is “stateless”

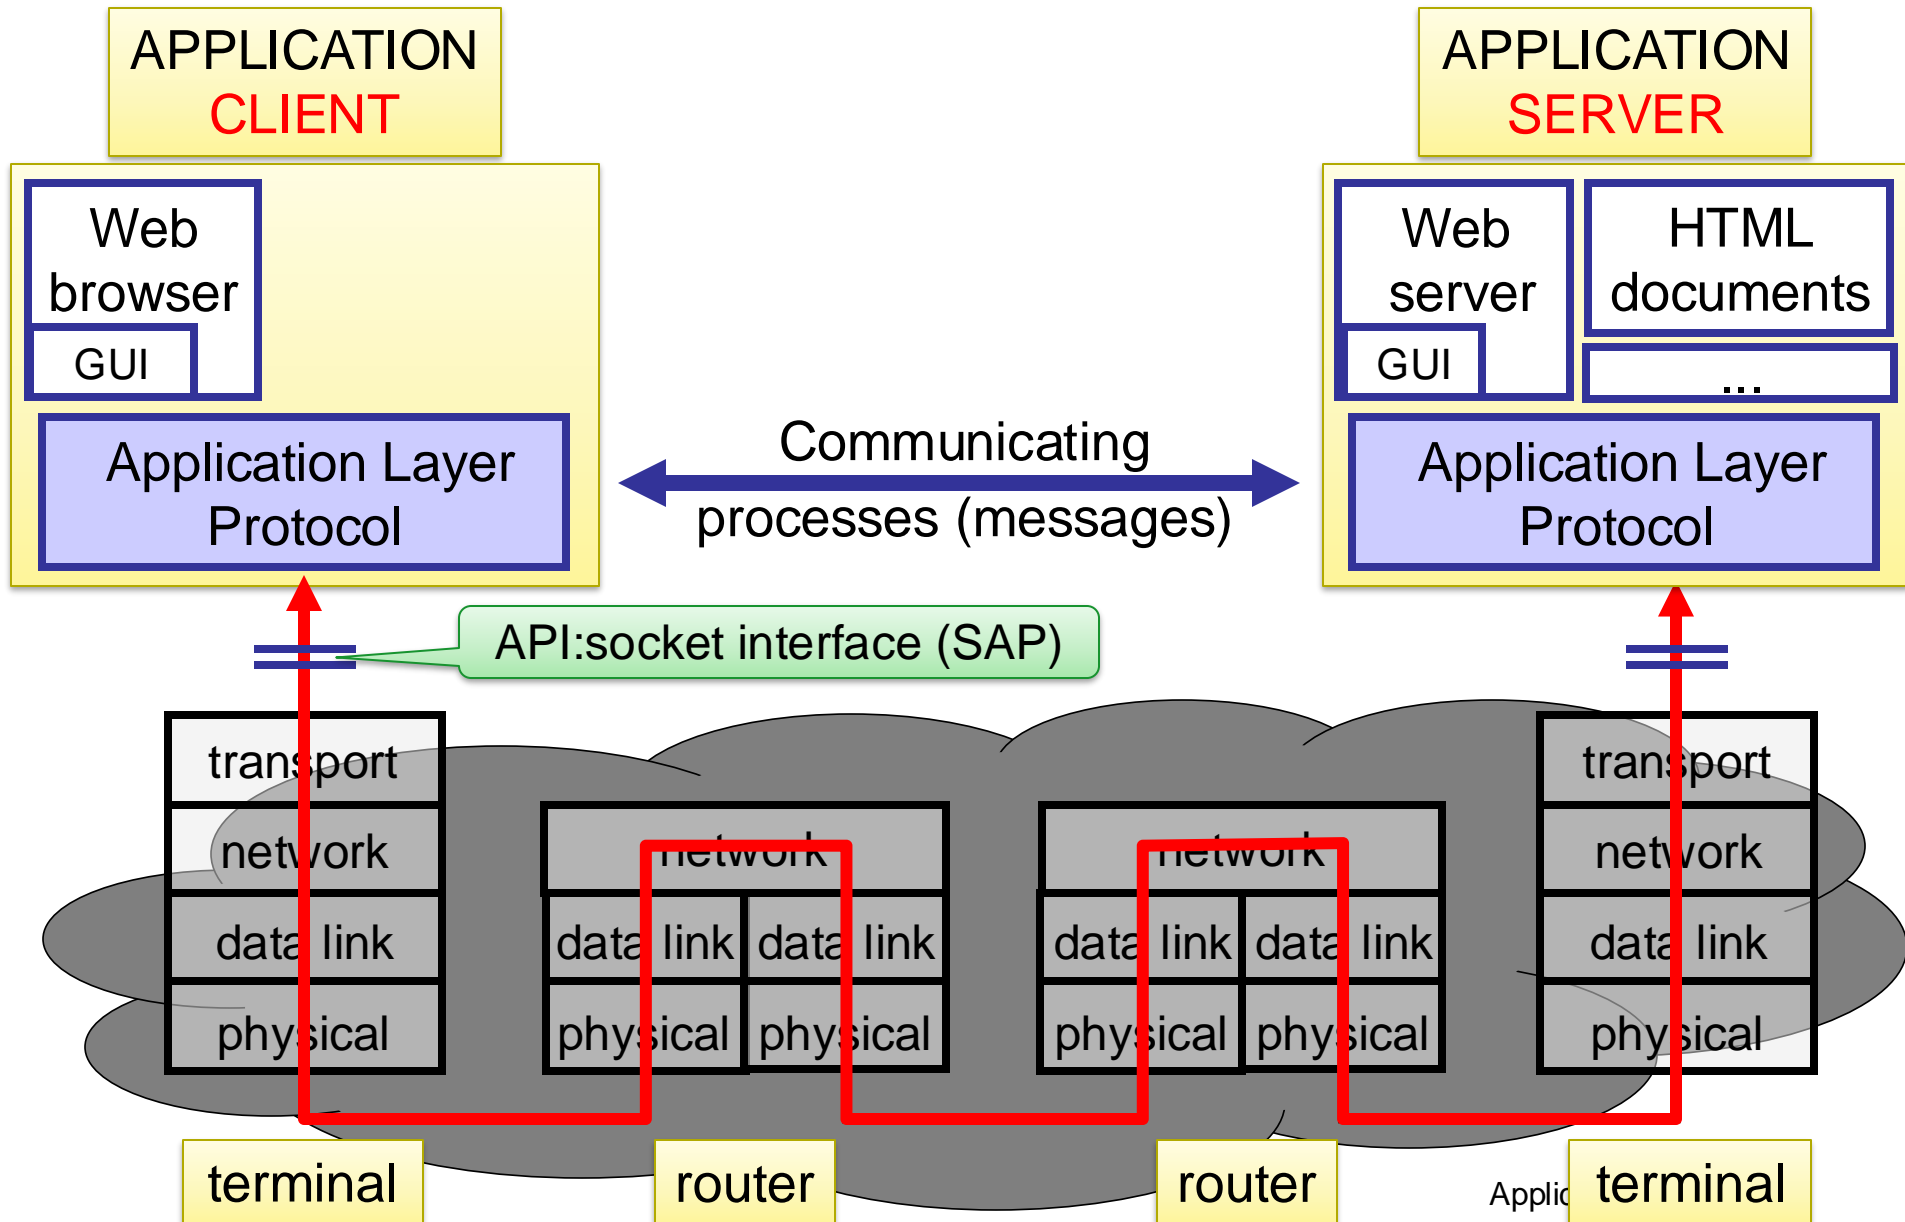
- server maintains no information about past client requests

aside

### protocols that maintain “state” are complex!

- past history (state) must be maintained
- if server/client crashes, their views of “state” may be inconsistent, must be reconciled

# HTTP in the global picture



# What a browser does in background ...

```
> telnet ugent.be 80
```

```
Trying 157.193.43.50...
```

```
Connected to ugent.be.
```

Set up a TCP connection

```
Escape character is '^['.
```

```
GET / HTTP/1.1
```

```
Host: ugent.be
```

Issue HTTP requests

```
HTTP/1.1 301 Moved Permanently
```

```
Date: Mon, 16 Sep 2024 14:06:01 GMT
```

```
Server: Apache
```

```
Location: https://ugent.be/
```

```
Content-Length: 225
```

```
Content-Type: text/html; charset=iso-8859-1
```

Receive HTTP response

```
<!DOCTYPE HTML PUBLIC "-//IETF//DTD HTML 2.0//EN">
```

```
<html><head>
```

```
<title>301 Moved Permanently</title>
```

```
</head><body>
```

```
<h1>Moved Permanently</h1>
```

```
<p>The document has moved <a href="https://ugent.be/">here</a>.</p>
```

```
</body></html>
```

```
Connection closed by foreign host.
```

Close TCP connection

Q: Try this at home ... why are all modern webserver returning 301 ?

# HTTP Connection Persistence

## non-persistent HTTP

- at most one object sent over TCP connection  
connection then closed
- downloading multiple objects requires multiple connections
- HTTP/1.0 uses non-persistent HTTP

## persistent HTTP

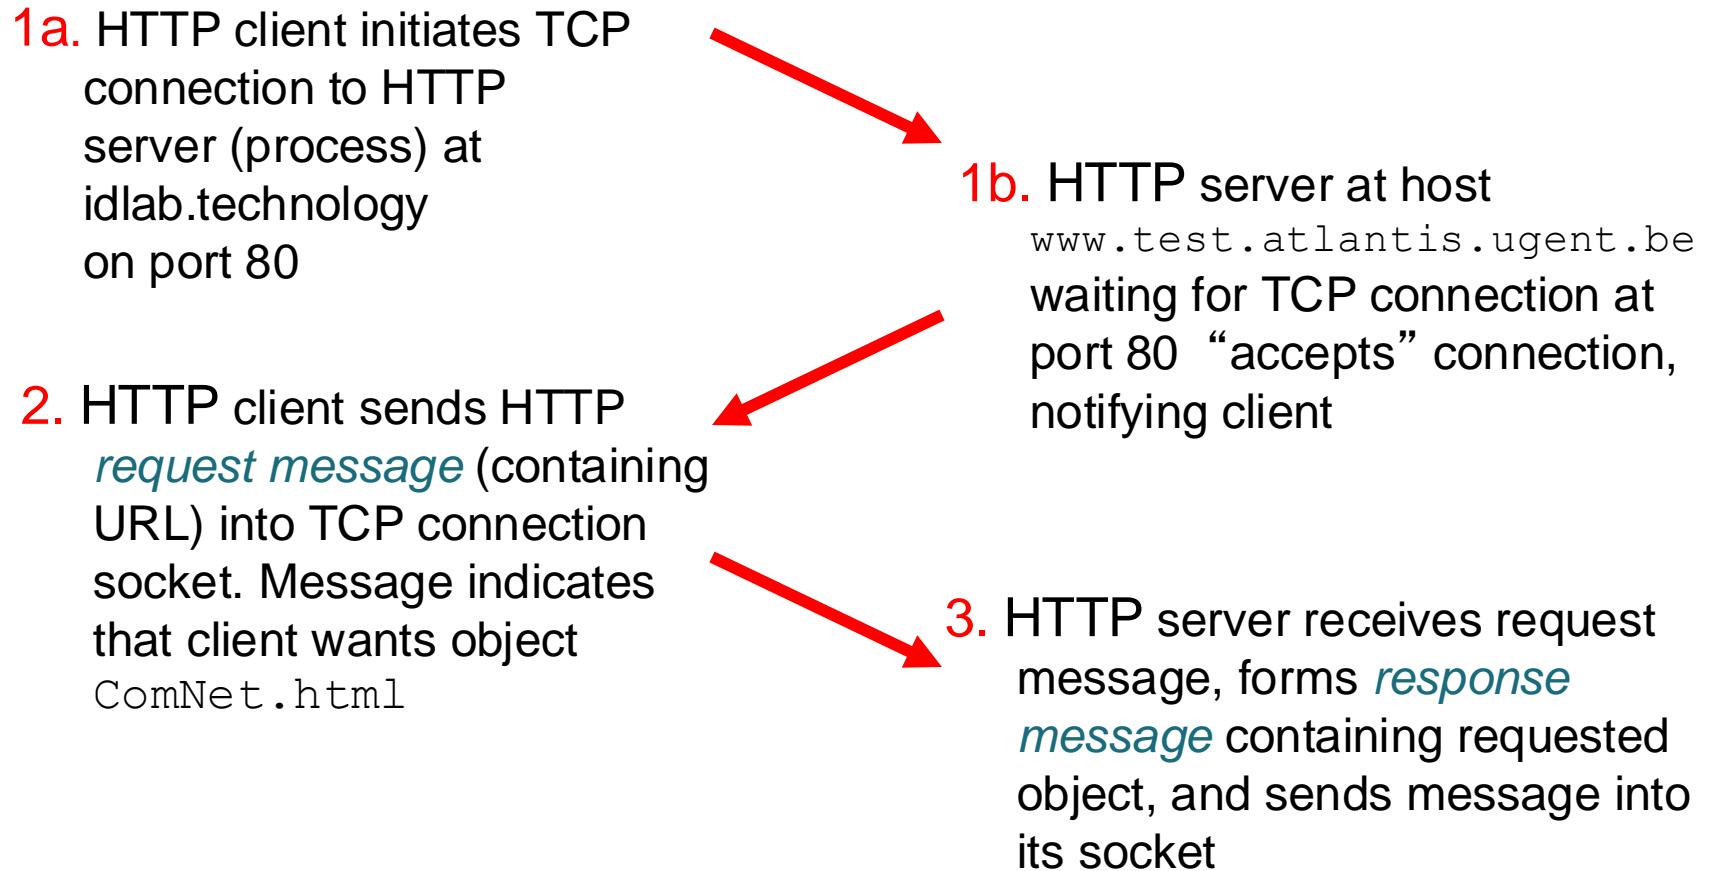
- multiple objects can be sent over single TCP connection between client, server
- HTTP/1.1 uses persistent connections in default mode

# Non-persistent HTTP

(contains text,  
references to 3  
jpeg images)

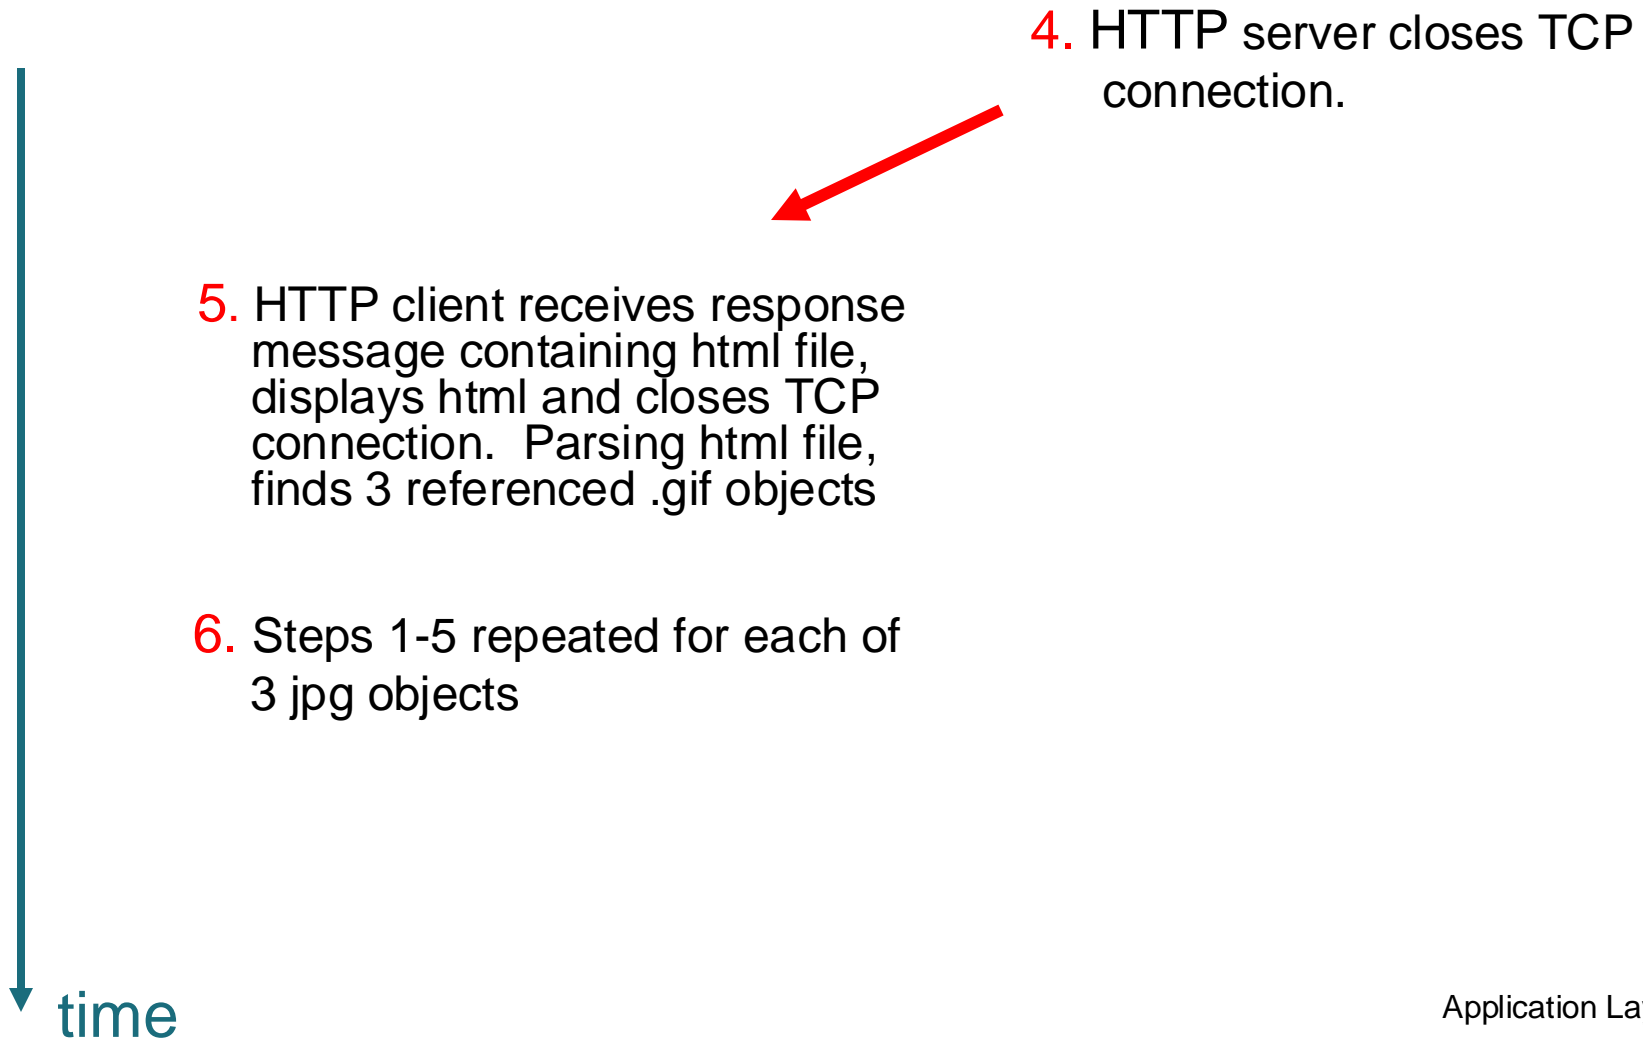
Suppose user enters in the browser the URL

`http://www.test.atlantis.ugent.be/ComNet.html`



time

# Non-persistent HTTP (cont.)



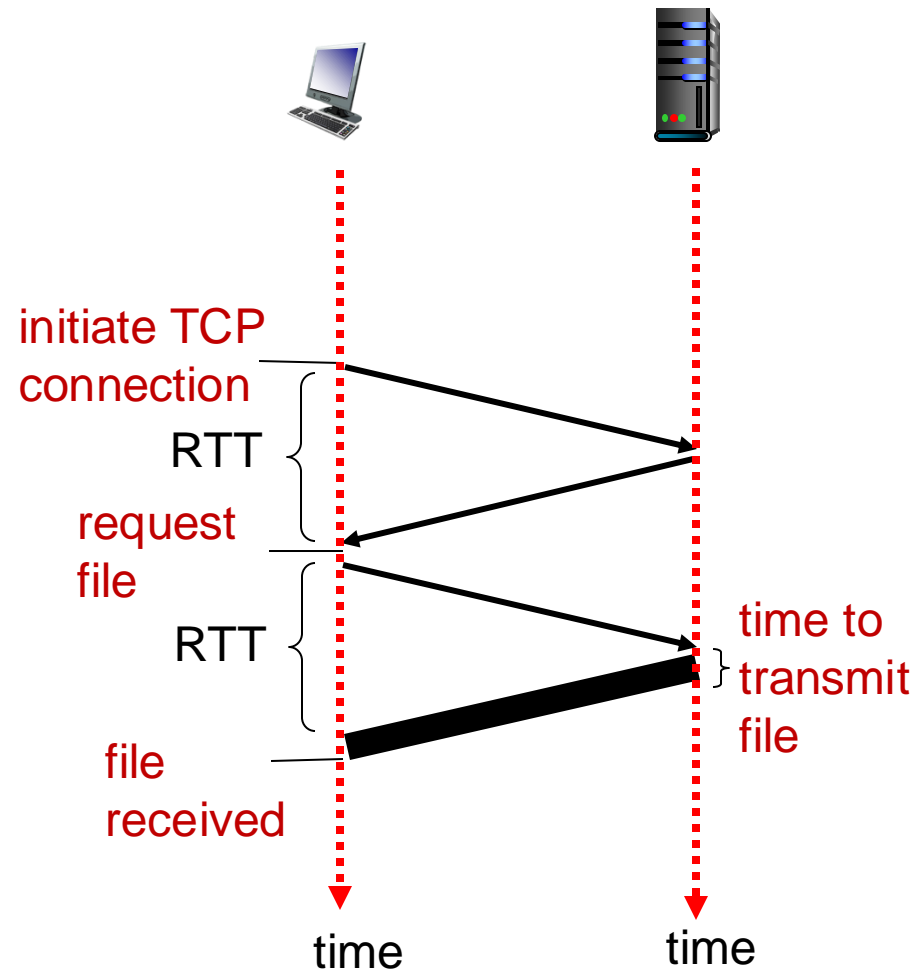
# Non-persistent HTTP: Response Time

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- file transmission time

total =  $2RTT + \text{file transmit time}$

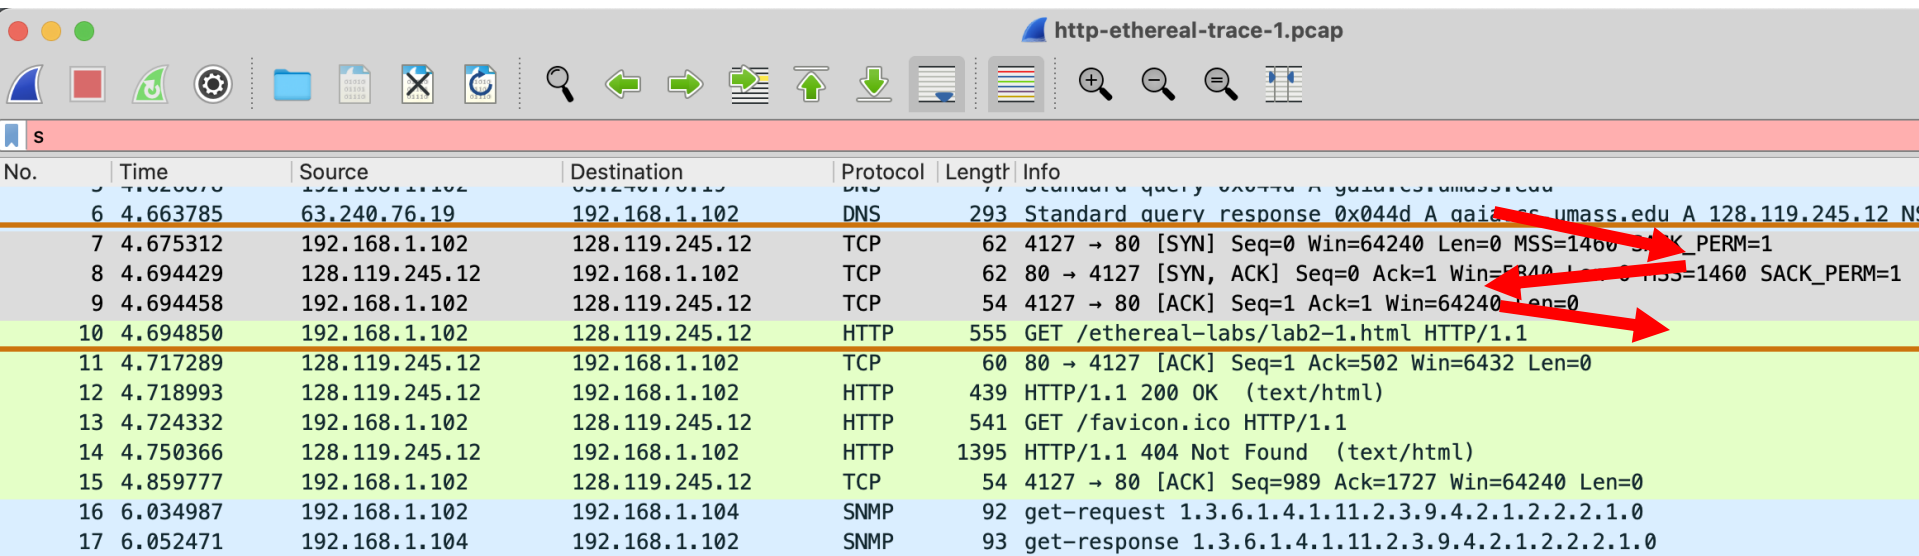
In practice:

1 KB @ 80 Kbps & 50 ms RTT  
 $(10\text{KBps} \times 1\text{KB}) + 2 \times 0,05\text{s} = 0,2\text{s}$   
Effective throughput = 40 Kbps



**Definition of RTT (Round Trip Time)** : time to send a small packet to travel from client to server and back

# TCP connection setup with webserver in Wireshark



http-ethereal-trace-1.pcap

No.	Time	Source	Destination	Protocol	Length	Info
6	4.663785	63.240.76.19	192.168.1.102	DNS	293	Standard query response 0x044d A gaias... umass.edu A 128.119.245.12 NS
7	4.675312	192.168.1.102	128.119.245.12	TCP	62	4127 → 80 [SYN] Seq=0 Win=64240 Len=0 MSS=1460 SACK_PERM=1
8	4.694429	128.119.245.12	192.168.1.102	TCP	62	80 → 4127 [SYN, ACK] Seq=0 Ack=1 Win=5840 Len=0 MSS=1460 SACK_PERM=1
9	4.694458	192.168.1.102	128.119.245.12	TCP	54	4127 → 80 [ACK] Seq=1 Ack=1 Win=64240 Len=0
10	4.694850	192.168.1.102	128.119.245.12	HTTP	555	GET /ethereal-labs/lab2-1.html HTTP/1.1
11	4.717289	128.119.245.12	192.168.1.102	TCP	60	80 → 4127 [ACK] Seq=1 Ack=502 Win=6432 Len=0
12	4.718993	128.119.245.12	192.168.1.102	HTTP	439	HTTP/1.1 200 OK (text/html)
13	4.724332	192.168.1.102	128.119.245.12	HTTP	541	GET /favicon.ico HTTP/1.1
14	4.750366	128.119.245.12	192.168.1.102	HTTP	1395	HTTP/1.1 404 Not Found (text/html)
15	4.859777	192.168.1.102	128.119.245.12	TCP	54	4127 → 80 [ACK] Seq=989 Ack=1727 Win=64240 Len=0
16	6.034987	192.168.1.102	192.168.1.104	SNMP	92	get-request 1.3.6.1.4.1.11.2.3.9.4.2.1.2.2.1.0
17	6.052471	192.168.1.104	192.168.1.102	SNMP	93	get-response 1.3.6.1.4.1.11.2.3.9.4.2.1.2.2.1.0



# Persistent HTTP

## Nonpersistent HTTP issues:

- requires 2 RTTs per object
- OS must work and allocate host resources for each TCP connection
- but browsers often open parallel TCP connections to fetch referenced objects

## Persistent HTTP

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server are sent over connection



## Persistent without pipelining:

- client issues new request only when previous response has been received
- one RTT for each referenced object

## Persistent with pipelining:

- default in HTTP/1.1
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects

# HTTP messages



## Request message

- Client -> Server
- Message structure

Method + url

headers

body

- **Request methods:**
  - GET
  - POST
  - HEAD
  - PUT
  - DELETE

## Response message

- Server -> Client
- Message structure

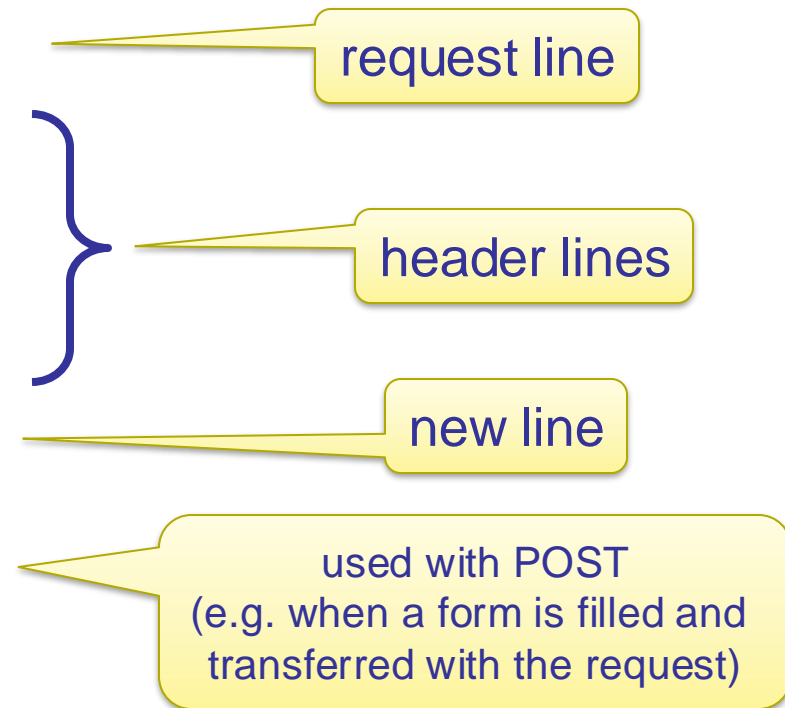

status

headers

body

- **Status codes:**
  - 200
  - 301
  - 400
  - 404
  - 505

# HTTP : Request Message

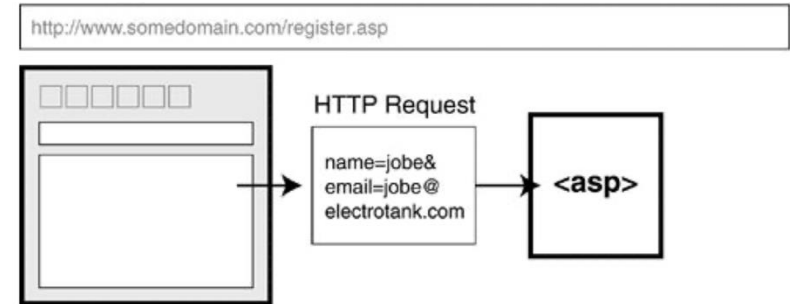


```
GET / HTTP/1.1                                     Request line
Host: neverssl.com
Connection: keep-alive
Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8
User-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X 10_15_7) AppleWebKit/605.1.15 (KHTML, like Gecko) Version/17.4 Safari/605.1.15
Accept-Language: en-GB,en;q=0.9                    Header lines
```

# Uploading Form Input


## POST method:

- web page often includes form input
- input is uploaded to server in entity body

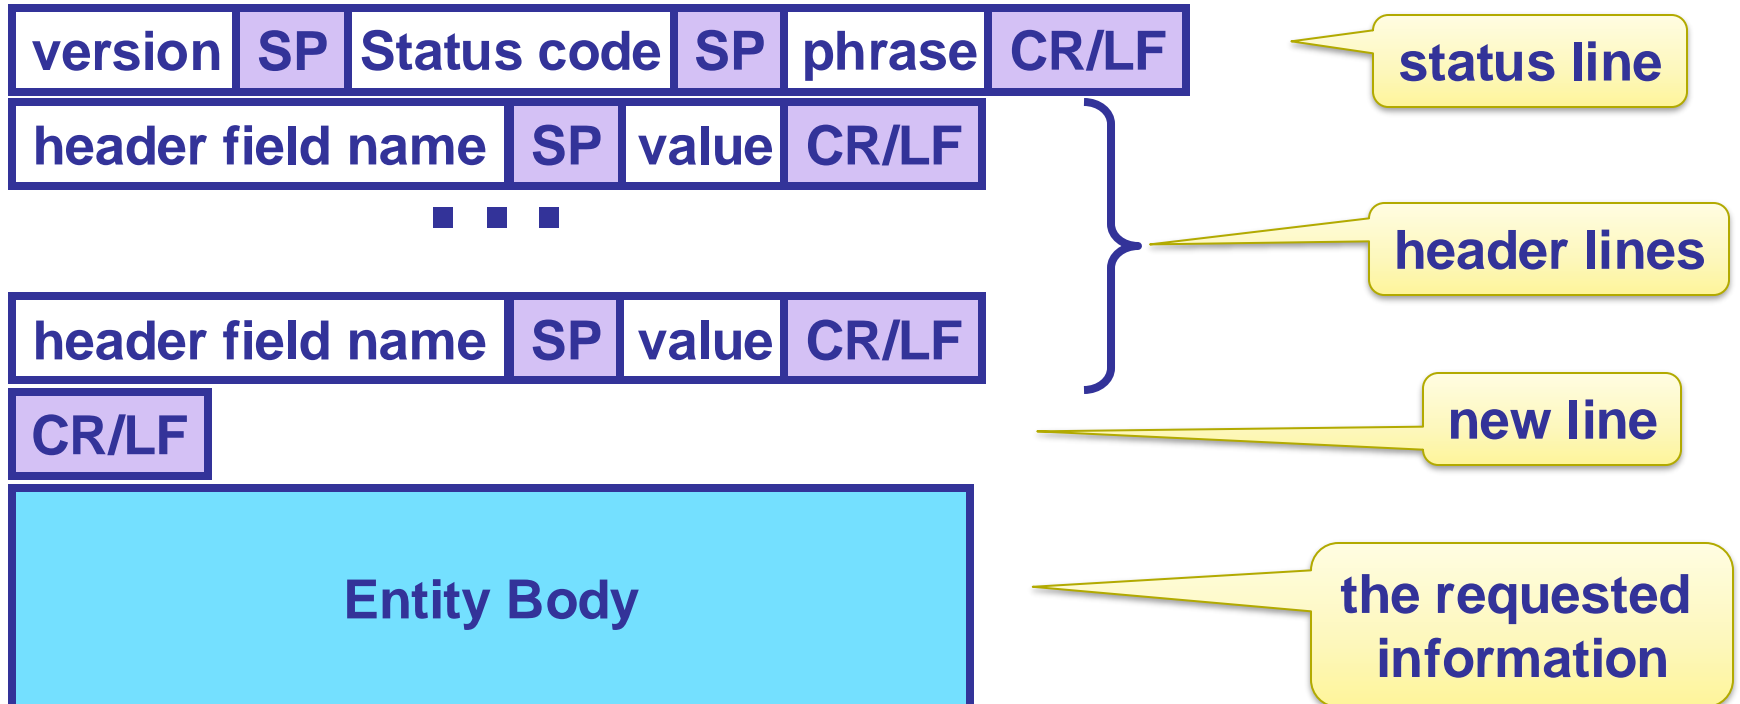


## URL method:

- uses GET method
- input is uploaded in URL field of request line:



# HTTP : Response Message



```
HTTP/1.1 200 OK
Date: Mon, 16 Sep 2024 14:32:25 GMT
Server: Apache/2.4.58 ()
Last-Modified: Wed, 29 Jun 2022 00:23:22 GMT
ETag: "8be-5e28b29291e10-gzip"
Accept-Ranges: bytes
Vary: Accept-Encoding
Content-Encoding: gzip
Content-Length: 1173
Keep-Alive: timeout=5, max=99
Connection: Keep-Alive
Content-Type: text/html; charset=UTF-8
```

Status line

Header lines

# HTTP Response Status

- Status code appears in 1st line in server-to-client response message.
- Most common = 200 – success
- Redirect
  - 301 permanent redirect
    - E.g. from http to https-page
  - 302 temporary
- 4xx error
  - Client did something wrong
- 5xx error
  - The server did something wrong

## 2xx Success

200

**Success / OK**

## 3xx Redirection

301

**Permanent Redirect**

302

**Temporary Redirect**

304

**Not Modified**

## 4xx Client Error

401

**Unauthorized Error**

403

**Forbidden**

404

**Not Found**

405

**Method Not Allowed**

## 5xx Server Error

501

**Not Implemented**

502

**Bad Gateway**

503

**Service Unavailable**

504

**Gateway Timeout**

# HTTP example (as in the old days)

```
> telnet neverssl.com 80
Trying 34.223.124.45...
Connected to neverssl.com.
Escape character is '^]'.
GET / HTTP/1.1
Host: neverssl.com
```

connection set-up

request to get  
homepage

```
HTTP/1.1 200 OK
Date: Mon, 16 Sep 2024 14:43:14 GMT
Server: Apache/2.4.58 ()
Upgrade: h2,h2c
Connection: Upgrade
Last-Modified: Wed, 29 Jun 2022 00:23:33 GMT
ETag: "f79-5e28b29d38e93"
Accept-Ranges: bytes
Content-Length: 3961
Vary: Accept-Encoding
Content-Type: text/html; charset=UTF-8
```

reply header

```
<html>
```

```
...
```

```
    <h2>How?</h2>
```

```
    <p>neverssl.com will never use SSL (also
known as TLS). No encryption, no strong
authentication, no HTTP/2.0, just plain old
unencrypted HTTP and forever stuck in the dark
ages of internet security.</p>
```

```
...
```

```
</html>
```

```
Connection closed by foreign host
```

HTML document  
neverssl homepage

close connection

# HTTP GET IN WIRESHARK

Wireshark File Edit View Go Capture Analyze Statistics Telephony Wireless Tools Help

USB 10/100/1000 LAN: en6

ip.addr == 185.47.29.183

No.	Time	Source	Destination	Protocol	Length	Info
3695	83.204758	192.168.1.151	185.47.29.183	TCP	78	55057 → 80 [SYN] Seq=0 Win=65535 Len=0 MSS=1460 WS=64 TSval=1083357673 TSecr=0 SACK_PERM=1
3701	83.223980	185.47.29.183	192.168.1.151	TCP	74	80 → 55057 [SYN, ACK] Seq=0 Ack=1 Win=8192 Len=0 MSS=1460 WS=256 SACK_PERM=1 TSval=304193469 TSecr=0
3702	83.224038	192.168.1.151	185.47.29.183	TCP	66	55057 → 80 [ACK] Seq=1 Ack=1 Win=131712 Len=0 TSval=1083357692 TSecr=304193469
3703	83.224620	192.168.1.151	185.47.29.183	HTTP	443	GET / HTTP/1.1
3704	83.309939	192.168.1.151	185.47.29.183	TCP	443	[TCP Retransmission] 55057 → 80 [PSH, ACK] Seq=1 Ack=1 Win=131712 Len=377 TSval=1083357777 TSecr=304193469
3705	83.342027	185.47.29.183	192.168.1.151	TCP	78	80 → 55057 [ACK] Seq=1 Ack=378 Win=66560 Len=0 TSval=304193481 TSecr=1083357692 SLE=1 SRE=378
3707	83.507590	185.47.29.183	192.168.1.151	TCP	1514	80 → 55057 [ACK] Seq=1 Ack=378 Win=66560 Len=1448 TSval=304193496 TSecr=1083357692 [TCP segment of data flow 0x12345678]
3708	83.507613	185.47.29.183	192.168.1.151	TCP	1514	80 → 55057 [ACK] Seq=1449 Ack=378 Win=66560 Len=1448 TSval=304193496 TSecr=1083357692 [TCP segment of data flow 0x12345678]
3709	83.507673	192.168.1.151	185.47.29.183	TCP	66	55057 → 80 [ACK] Seq=378 Ack=2897 Win=128832 Len=0 TSval=1083357973 TSecr=304193496
3713	83.527867	185.47.29.183	192.168.1.151	TCP	1514	80 → 55057 [ACK] Seq=2897 Ack=378 Win=66560 Len=1448 TSval=304193499 TSecr=1083357973 [TCP segment of data flow 0x12345678]
3714	83.527868	185.47.29.183	192.168.1.151	HTTP	828	HTTP/1.1 200 OK (text/html)
3715	83.527930	192.168.1.151	185.47.29.183	TCP	66	55057 → 80 [ACK] Seq=378 Ack=5107 Win=128832 Len=0 TSval=1083357992 TSecr=304193499
3747	83.575273	192.168.1.151	185.47.29.183	HTTP	486	GET /tmp/cache/stylesheet_combined_e68d8027adb7e846a313e5e6f01b4e74.css HTTP/1.1
3749	83.576143	192.168.1.151	185.47.29.183	TCP	78	55058 → 80 [SYN] Seq=0 Win=65535 Len=0 MSS=1460 WS=64 TSval=1083358034 TSecr=0 SACK_PERM=1
3750	83.576547	192.168.1.151	185.47.29.183	TCP	78	55059 → 80 [SYN] Seq=0 Win=65535 Len=0 MSS=1460 WS=64 TSval=1083358035 TSecr=0 SACK_PERM=1

▶ Frame 3703: 443 bytes on wire (3544 bits), 443 bytes captured (3544 bits) on interface 0

▶ Ethernet II, Src: 00:e0:4c:68:00:45, Dst: 24:f5:a2:ba:9b:ab

▶ Internet Protocol Version 4, Src: 192.168.1.151, Dst: 185.47.29.183

▶ Transmission Control Protocol, Src Port: 55057, Dst Port: 80, Seq: 1, Ack: 1, Len: 377

▼ Hypertext Transfer Protocol

▶ GET / HTTP/1.1\r\n

Host: www.hortamuseum.be\r\n

User-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X 10.15; rv:79.0) Gecko/20100101 Firefox/79.0\r\n

Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/webp,\*/\*;q=0.8\r\n

Accept-Language: en-US,en;q=0.5\r\n

Accept-Encoding: gzip, deflate\r\n

Referer: https://www.google.com/\r\n

Connection: keep-alive\r\n

Upgrade-Insecure-Requests: 1\r\n\r\n

[Full request URI: <http://www.hortamuseum.be/>]

[HTTP request 1/6]

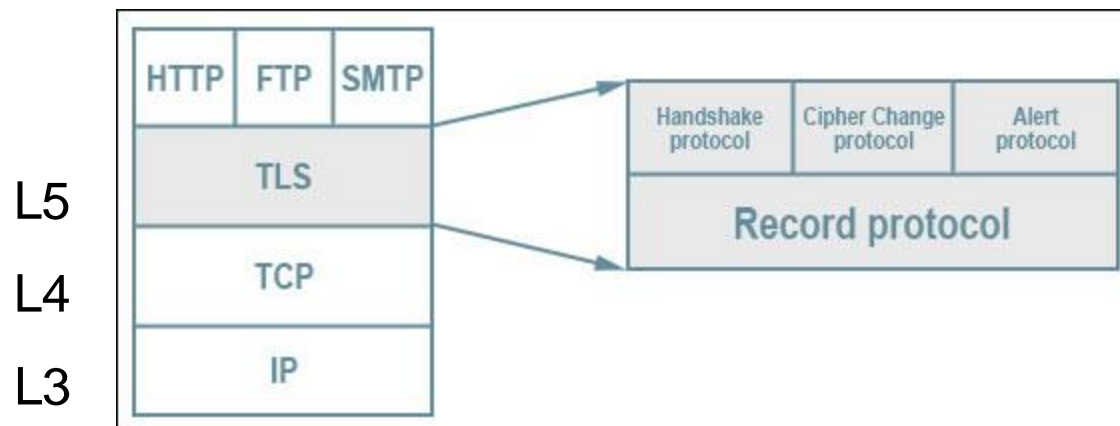
[Response in frame: 3714]

[Next request in frame: 3747]



# Transport Layer Security (TLS)

- (plain) HTTP relies directly on TCP
  - All application layer communication is sent unencrypted over TCP over the Internet
  - “anyone” can monitor what web pages you visit, etc.
- Alternative = Transport Layer Security (also Secure Sockets Layer):
  - Set up an encrypted connection between client/server first over TCP -> Cfr. Chapter 8
  - Exchange all application layer traffic (HTTP) over the encrypted channel -> HTTPS



# HTTP today (using HTTPS)

```
$ openssl s_client -crlf -connect www.ugent.be:443
```

```
<certification handshake>
```

```
GET / HTTP/1.1
```

```
HTTP/1.1 200 OK
```

```
Date: Thu, 23 Sep 2021 09:35:38 GMT
```

```
Server: waitress
```

```
...
```

```
Accept-Ranges: bytes
```

```
Via: 1.1 www.ugent.be
```

```
Set-Cookie: ROUTEID=.1; path=/
```

```
<html>
```

```
  <head>
```

```
    <title>UNIVERSITEIT GENT - UNIVERSITY OF  
    GHENT</title>
```

```
    <style type="text/css">
```

```
      ...
```

```
    </style>
```

```
</head>
```

```
<body bgcolor="#000066" link="#cccccc"  
  vlink="#cccccc" alink="#666666">
```

```
  ...
```

```
</body>
```

```
</html>
```

```
Connection closed by foreign host
```

}  
secure  
connection set-up  
using TLS Chapter 8  
}

# Web session to ugent.be -> HTTPS

USB 10/100/1000 LAN: en6

ip.addr == 157.193.43.50

No.	Time	Source	Destination	Protocol	Length	Info
193	4.789498	157.193.43.50	192.168.1.151	TLSv1.2	1514	Certificate [TCP segment of ...]
194	4.790555	157.193.43.50	192.168.1.151	TLSv1.2	233	Server Key Exchange, Server ...
195	4.790601	192.168.1.151	157.193.43.50	TCP	66	50847 → 443 [ACK] Seq=518 Ac...
197	4.801573	192.168.1.151	157.193.43.50	TLSv1.2	159	Client Key Exchange, Change ...
198	4.815677	157.193.43.50	192.168.1.151	TLSv1.2	117	Change Cipher Spec, Encrypte...
199	4.815729	192.168.1.151	157.193.43.50	TCP	66	50847 → 443 [ACK] Seq=611 Ac...
200	4.816852	192.168.1.151	157.193.43.50	TLSv1.2	977	Application Data
212	4.838329	157.193.43.50	192.168.1.151	TLSv1.2	594	Application Data
213	4.838385	192.168.1.151	157.193.43.50	TCP	66	50847 → 443 [ACK] Seq=1522 A...
215	4.969937	192.168.1.151	157.193.43.50	TLSv1.2	971	Application Data
216	4.983714	157.193.43.50	192.168.1.151	TLSv1.2	534	Application Data
217	4.983782	192.168.1.151	157.193.43.50	TCP	66	50845 → 443 [ACK] Seq=2382 A...
218	5.001301	192.168.1.151	157.193.43.50	TLSv1.2	958	Application Data
219	5.021851	157.193.43.50	192.168.1.151	TLSv1.2	570	Application Data
220	5.021903	192.168.1.151	157.193.43.50	TCP	66	50845 → 443 [ACK] Seq=3274 A...

▶ Frame 200: 977 bytes on wire (7816 bits), 977 bytes captured (7816 bits) on interface 0

▶ Ethernet II, Src: 00:e0:4c:68:00:45, Dst: 24:f5:a2:ba:9b:ab

▶ Internet Protocol Version 4, Src: 192.168.1.151, Dst: 157.193.43.50

▶ Transmission Control Protocol, Src Port: 50847, Dst Port: 443, Seq: 611, Ack: 5763, Len: 911

▼ Secure Sockets Layer

▼ TLSv1.2 Record Layer: Application Data Protocol: http-over-tls

Content Type: Application Data (23)

Version: TLS 1.2 (0x0303)

Length: 906

Encrypted Application Data: 0000000000000017c50414cc4a404a9511da72eefc190db...

L2  
L3  
L4  
L5

THE WORLD'S CITY.  
IT'S WHEREVER  
YOU ARE.

# The New York Times

Monday, October 14, 2013 Last Update: 1:42 PM ET

citi  
The World's Citi™

 Search

Follow Us Personalize Your Weather

Why is the homepage of my newspaper “aware” of my interest in bank loans and intentions to buy a house?

WHEREVER THERE'S  
OPPORTUNITY,  
WE'RE THERE TO HELP  
MAKE IT REAL.

citi  
The World's Citi™

Colleges Ad

WORLD  
U.S.  
POLITICS  
NEW YORK  
BUSINESS  
DEALBOOK  
TECHNOLOGY  
SPORTS  
SCIENCE  
HEALTH  
ARTS  
STYLE  
OPINION

Autos  
Blogs

## FISCAL CRISIS

### Seeking Deal to Avert Default, Lawmakers to Meet Obama

By JONATHAN WEISMAN  
12:51 PM ET

Senate talks on reopening the government and raising the debt ceiling were expected to accelerate Monday afternoon, and President Obama was to meet with top Congressional leaders from both parties.



RETRO REPORT | VIDEO

### Portents of Doom in a Sheep Named Dolly

By NICHOLAS WATTS

## The Opinion Pages

How the G.O.P. Is Shaping the Midterms  
Republicans have alienated the public. In 2014, can independents capitalize on that?



### MORE IN OPINION

- Editorial: False Equality in Michigan
- Stiglitz: Inequality Is a Choice
- Op-Ed: The Bay of Bengal

TURNING THE PAGE  
How We Called It, Down Through the Years  
A look back at a selection of columns, essays and editorials that were published in The International Herald Tribune.  
• Visit the Special Section »

### OP-ED COLUMNISTS

- Keller: Obamacare, the Rest of the Story
- Krugman: The Dixiecrat Solution

# User-Server State: Cookies

Many Web sites use cookies

## Cookies rely on four components:

- 1) cookie header line of HTTP **response** message
- 2) cookie header line in next HTTP **request** message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

## example:


- Susan always accesses Internet from PC
- visits specific e-commerce site for first time
- when initial HTTP requests arrives at site, site creates:
  - unique ID
  - entry in backend database for ID

# Cookies: keeping “state”

client



server



cookie file

usual http request msg

Amazon server  
creates ID  
1678 for user

usual http response  
set-cookie: 1678



ebay 8734  
amazon 1678

create  
entry


backend  
database

usual http request msg  
cookie: 1678

cookie-  
specific  
action

access

usual http response msg



access

cookie-  
specific  
action

one week later:



ebay 8734  
amazon 1678

usual http request msg  
cookie: 1678

usual http response msg

# Cookies

## *what cookies can be used for:*


- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)

## *how to keep “state”:*

- protocol endpoints: maintain state at sender/receiver over multiple transactions
- cookies: http messages carry state

### *cookies and privacy:* aside

- cookies permit sites to learn a lot about you
- you may supply name and e-mail to sites
- EU GDPR law requires explicit agreement for use of cookies since 2018



### Privacy Policy

This website uses cookies to ensure you get the best experience on our website.


[Read More](#)

[Got it!](#)

# Conditional GET: client-side caching



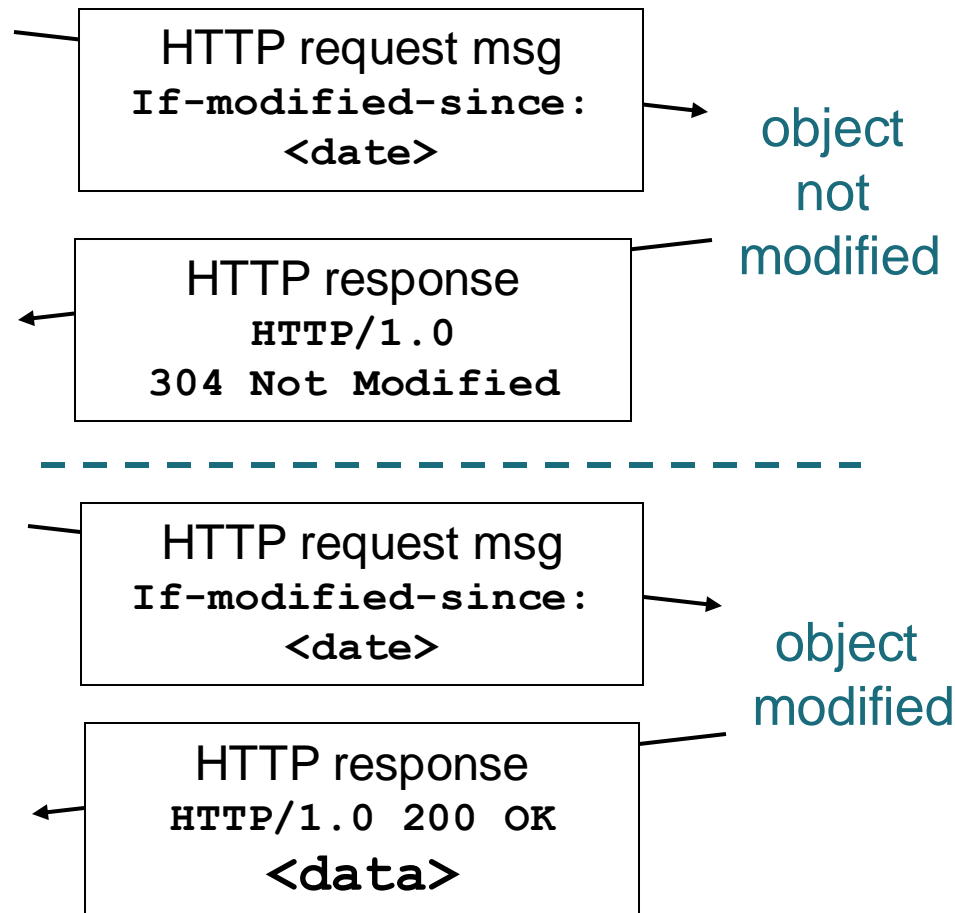
client



server

**Goal:** don't send object if client has up-to-date cached version

- no object transmission delay
- lower link utilization
- client: specify date of cached copy in HTTP request  
`If-modified-since: <date>`
- server: response contains no object if cached copy is up-to-date:  
`HTTP/1.0 304 Not Modified`






# Web caches (proxy server)

**Goal:** satisfy client request without involving origin server

- user sets browser: Web accesses via cache
- browser sends all HTTP requests to cache
  - object in cache: cache returns object
  - else cache requests object from origin server, then returns object to client



# More About Web Caching

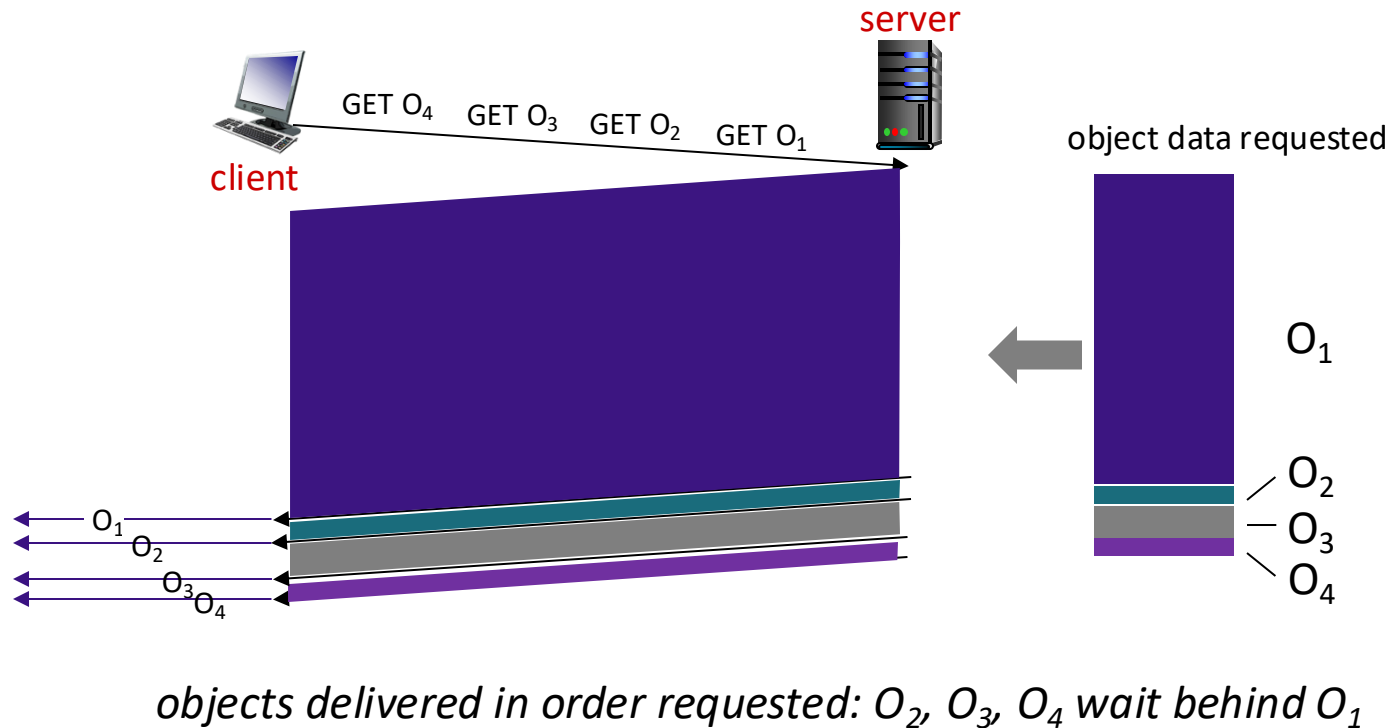
- cache acts as both client and server
  - server for original requesting client
  - client to origin server
- typically cache is installed by ISP (university, company, residential ISP)

## **why Web caching?**

- reduce response time for client request
- reduce traffic on an institution's access link
- Internet dense with caches: enables “poor” content providers to effectively deliver content (so too does P2P file sharing)

# HTTP 1.1 Head of Line Blocking

HTTP 1.1: client requests 1 large object (e.g., video file) and 3 smaller objects



# HTTP/2

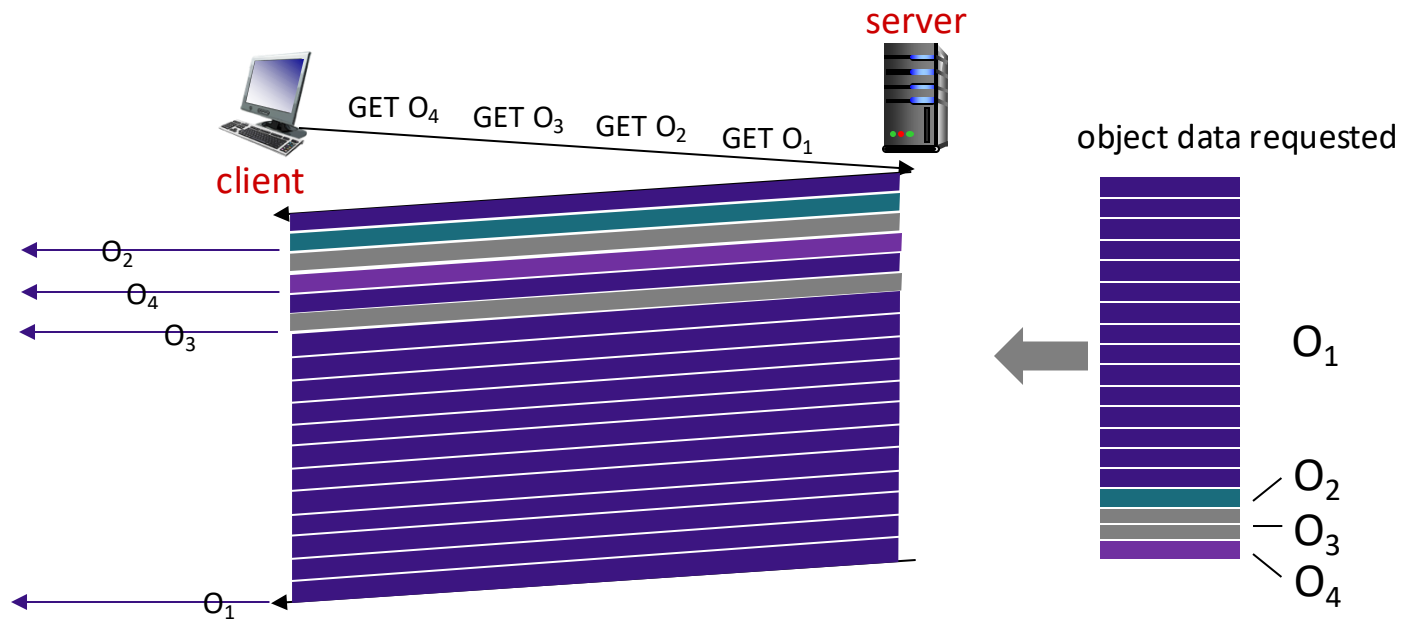
*Key goal:* decreased delay in multi-object HTTP requests

HTTP/2: [RFC 7540, 2015] increased flexibility at *server* in sending objects to client:

- methods, status codes, most header fields unchanged from HTTP 1.1
- transmission order of requested objects based on client-specified object priority (not necessarily FCFS)
- *push* unrequested objects to client
- divide objects into frames, schedule frames to mitigate HOL blocking

# HTTP/2: mitigating HOL blocking

HTTP/2: objects divided into frames, frame transmission interleaved



*O<sub>2</sub>, O<sub>3</sub>, O<sub>4</sub> delivered quickly, O<sub>1</sub> slightly delayed*

# HTTP evolution and versions

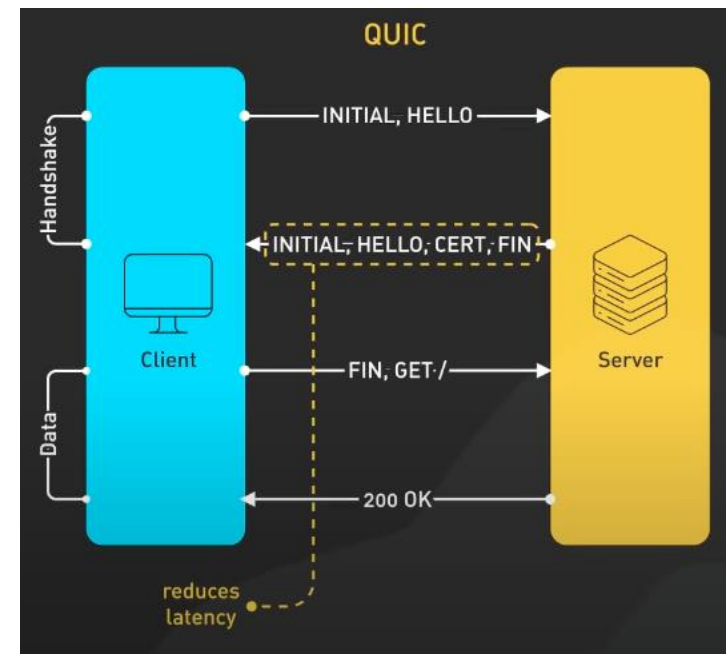
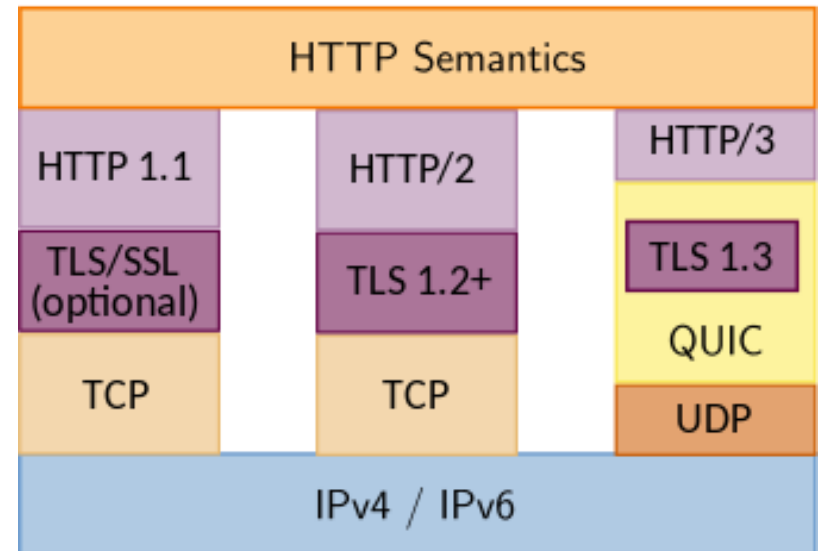


HTTP/1.0	HTTP/1.1	HTTP2.0
GET, POST	Persistent connections	Improve page load speed (decreasing latency) <ul style="list-style-type: none"><li>• Data compression of HTTP headers</li><li>• Multiplexing multiple requests</li><li>• HTTP/2 Server Push</li></ul>
HEAD asks server to leave requested object out of response	OPTIONS method Used to determine abilities of the server.	Longer lived (~permanent) HTTP connections
PUT uploads file in entity body to path specified in URL field DELETE deletes file specified in the URL field	Additional caching, authentication & compression options	HTTP/2 is a binary protocol instead of text-based (as 1.*)

# HTTP/2 to HTTP/3

HTTP/2 over single TCP connection means:


- recovery from packet loss still stalls all object transmissions
  - as in HTTP 1.1, browsers have incentive to open multiple parallel TCP connections to reduce stalling, increase overall throughput
- no security over vanilla TCP connection
- HTTP/3: adds security, per object error- and congestion-control (more pipelining) over UDP
  - more on HTTP/3 in transport layer



[Overview](#)  
[On Youtube](#)



# Adoption



Source: ByteByteGo (2024)



# Experiment in Kathará



- Chapter 2 -> http experiment
  - Involves commonly used Apache 2 webserver(s)
  - Use wget to fetch page including images
- Fetch homepage of server on client
  - Which HTTP version?
  - Which HTTP messages?
  - How many TCP connections?

